GPU CLOUD, LLM HARDWARE & AI INFRASTRUCTURE

COMPREHENSIVE COMPARISON - INDIA 2026

March 2026 | 1 USD = Rs.86.5 | All prices in Indian Rupees (INR)

DOCUMENT SCOPE This report covers all LLM inference hardware and GPU cloud platforms relevant to Indian AI developers as of March 2026 - from desktop mini-PCs to cloud GPU clusters to edge devices. Metrics include: token generation speed (tok/s), prompt evaluation rate, power consumption, cooling requirements, GPU utilisation, multi-user serving capacity, cost in INR, data sovereignty, and India compliance.

8 Hardware Platforms | 9 Global Cloud Providers | 7 India DC/Cloud Providers

1. Hardware Devices - Specifications & Performance

Compares all on-premise and edge hardware for running LLMs locally. Platforms: Mac Mini M4, Mac Studio, NVIDIA DGX Spark, NVIDIA Jetson AGX Thor, TAALAS 10x RTX 4090 cluster, Cloud H100, and Raspberry Pi 5 AI Kit.

1.1 Hardware Specifications Overview

Platform / Device	GPU / Chip	VRAM / RAM	Mem BW	Form Factor	Max LLM Size	Price (INR)	Category	CUDA
Mac Mini M4 (16 GB)	Apple M4 / 10-core GPU	16 GB unified	**120 GB/s**	Desktop mini	**7B INT4**	**Rs.57,900**	Personal AI	✗
Mac Mini M4 Pro (64 GB)	Apple M4 Pro / 20-core GPU	64 GB unified	**273 GB/s**	Desktop mini	**30B INT4**	**Rs.1,69,900**	Dev workstation	✗
Mac Studio M4 Max (128 GB)	Apple M4 Max / 40-core GPU	128 GB unified	**546 GB/s**	Desktop compact	**70B INT4 / 120B FP4**	**~Rs.4,19,000**	Power workstation	✗
NVIDIA DGX Spark	GB10 Grace Blackwell / 1 PFLOP FP4	128 GB LPDDR5X	**273 GB/s**	Desktop mini (gold)	**200B FP4 (2x linked)**	**Rs.4,06,454**	AI supercomputer	✓
NVIDIA Jetson AGX Thor	Blackwell GPU / 2070 FP4 TFLOPS	128 GB unified	**273 GB/s**	Edge / Embedded module	**70B FP8 / 120B FP4**	**~Rs.2,16,250**	Edge AI / Robotics	✓
TAALAS - 10x RTX 4090	10x NVIDIA RTX 4090 / 24 GB each	240 GB GDDR6X	**10+ TB/s aggregate**	Server rack 4U+	**70-130B INT4**	**~Rs.38-43 lakhs**	Production cluster	✓
Cloud H100 80GB (1x GPU)	NVIDIA H100 SXM 80GB HBM3	80 GB HBM3	**3.35 TB/s**	Cloud VM (no hardware)	**70B FP16 (full precision)**	**Rs.162-604/hr**	Pay-as-you-go	✓
Raspberry Pi 5 + AI Kit	BCM2712 + Hailo-8L 13 TOPS NPU	8 GB LPDDR4X	**~13 GB/s**	SBC 85x56mm	**1-3B (edge only)**	**~Rs.7,200**	Education/Prototype	✗

1.2 LLM Performance Benchmarks (Llama 3 / DeepSeek R1, Ollama / vLLM)

Benchmark note: Token gen = decode output speed (tok/s). Prompt eval = prefill speed. 7B column = Llama 3.1 7B Q4_K_M. 70B column = Llama 3.3 70B / DeepSeek R1 70B at INT4/Q4. * = 1-3B model only.

Platform	Tok/s 7B Q4	Tok/s 70B Q4	Prompt/s 7B	Prompt/s 70B	Power (W)	Cooling	GPU Util	Multi-user	Quality	Data Sovereignty
Mac Mini M4 (16 GB)	55	N/A	1,200	N/A	35 W	Passive + micro-fan. Silent.	60-80%	1 user	★★★☆☆	On-prem (full)
Mac Mini M4 Pro (64 GB)	95	6	1,800	140	55 W	Single fan. Very quiet.	70-90%	2-3 users	★★★★☆	On-prem (full)
Mac Studio M4 Max (128 GB)	180	28	3,200	420	100 W	Dual fan. Near-silent.	75-95%	3-5 users	★★★★☆	On-prem (full)
NVIDIA DGX Spark	220	38	4,500	520	170 W	Active fan. Moderate noise.	85-99%	3-8 users	★★★★★	On-prem (full)
NVIDIA Jetson AGX Thor	160	18	2,600	280	80 W	Compact active. Low noise.	80-95%	8-16 (small models)	★★★★☆	On-prem (full)
TAALAS - 10x RTX 4090	600	65	20,000	2,100	3,000 W	MUST HAVE active/liquid cooling. Loud.	80-95%	15-25 users	★★★★☆	On-prem (full)
Cloud H100 80GB (1x GPU)	1,000	105	28,000	2,800	700 W (DC managed)	Datacenter managed.	50-90%	8-20+ (scalable)	★★★★★	Region-dependent
Raspberry Pi 5 + AI Kit	8*	N/A	3 (1B)	N/A	12 W	Heatsink. Fanless optional.	N/A (NPU)	Demo only	★★☆☆☆	On-prem (full)

Power Efficiency Winner: Mac Mini M4 Pro = ~1.7 tok/W for 70B INT4. TAALAS is fastest absolute but 0.022 tok/W. Jetson Thor is best edge at ~0.22 tok/W with full CUDA+vLLM.

2. Global Cloud Platforms - Pricing, Comparison & Break-Even Analysis

Monthly costs at 8 hours/day usage. Break-even = months until cumulative cloud spending exceeds NVIDIA DGX Spark one-time cost (Rs.4,06,454). All on-demand rates, March 2026.

2.1 Global Cloud Platform Comparison (vs DGX Spark baseline)

Platform	GPU Hardware	RAM / vCPUs	USD/hr	INR/hr	Monthly INR (8h/day)	Data Sovereignty	Quality	Efficiency	Break-even	Verdict
AWS EC2 p5.48xlarge	H100 SXM 80GB (8x)	2 TB / 192 vCPU	**$3.93/GPU/hr**	**Rs.340/hr**	**Rs.2,44,795**	Mumbai avail. US-controlled	★★★★★	★★★★☆	**1.7**	Best ecosystem. SageMaker.
Google Cloud A3-High	H100 80GB	312 GB / 12 vCPU	**$3.00/GPU/hr**	**Rs.260/hr**	**Rs.1,86,840**	Mumbai avail. US-controlled	★★★★★	★★★★★	**2.2**	Best value hyperscaler. Vertex AI.
Azure NC H100 v5	H100 80GB	220 GB / 6 vCPU	**$6.98/GPU/hr**	**Rs.604/hr**	**Rs.4,34,749**	Limited India. MS-controlled	★★★★☆	★★★☆☆	**0.9**	OpenAI integration. Most expensive.
Oracle Cloud (Bare Metal)	H100 SXM 80GB (8x)	2 TB / 128 vCPU	**$1.25/GPU/hr**	**Rs.108/hr**	**Rs.77,850**	Hyderabad avail.	★★★★☆	★★★★☆	**5.2**	Cheapest hyperscaler per GPU.
Lambda Labs	H100 SXM 80GB	200 GB / 26 vCPU	**$2.99/GPU/hr**	**Rs.259/hr**	**Rs.1,86,235**	US/EU only - no India DC	★★★★☆	★★★★☆	**2.2**	No egress fees. ML-focused.
RunPod (on-demand)	H100 80GB	96 GB / 12 vCPU	**$1.99/GPU/hr**	**Rs.172/hr**	**Rs.1,23,955**	US/EU only - no India DC	★★★☆☆	★★★★☆	**3.3**	Budget dev/inference. Not prod.
CoreWeave	H100 SXM 80GB InfiniBand	192 GB / 12 vCPU	**$6.16/GPU/hr**	**Rs.533/hr**	**Rs.3,83,628**	US/EU only	★★★★★	★★★★★	**1.1**	Best distributed training.
Vast.ai (marketplace)	H100 80GB (spot)	64 GB / 8 vCPU	**$1.87/GPU/hr**	**Rs.162/hr**	**Rs.1,16,429**	Marketplace - community nodes	★★★☆☆	★★★☆☆	**3.5**	Absolute cheapest. Not for prod.
>> DGX Spark (BASELINE)	Blackwell GB10 / 1 PFLOP FP4	128 GB / 20 ARM cores	**--**	**--**	**Rs.4,06,454 ONE-TIME**	100% on-premises	★★★★☆	★★★★☆	**BASE**	Pays off vs Azure in < 30 days.

2.2 Break-Even Analysis - DGX Spark Rs.4,06,454 one-time vs Cloud Monthly

How to read: Shorter break-even = cloud MORE expensive per month. 0.9 months = DGX Spark pays for itself in under 30 days vs that provider.

Cloud Platform	INR/hr	Monthly INR (8h/day)	Break-Even vs DGX Spark	Notes & Recommendation
Azure NC H100 v5	Rs.604/hr	Rs.4,34,749	0.9 months	DGX Spark pays off in <30 days! Azure is the most expensive hyperscaler.
CoreWeave H100 InfiniBand	Rs.533/hr	Rs.3,83,628	1.1 months	Only worth it for large distributed training where InfiniBand is critical.
AWS EC2 p5.48xlarge	Rs.340/hr	Rs.2,44,795	1.7 months	DGX Spark cheaper after ~52 days at 8hrs/day. Mumbai DC good for India.
Google Cloud A3-High	Rs.260/hr	Rs.1,86,840	2.2 months	Best value hyperscaler. TPU alternative. Sustained-use discounts help.
Lambda Labs H100	Rs.259/hr	Rs.1,86,235	2.2 months	No egress fees. Clean billing. Good until month 2 of continuous use.
Yotta Shakti (India)	Rs.200/hr	Rs.1,44,000	2.8 months	Best Indian sovereign option. IndiaAI subsidy Rs.115-150/hr extends to 4.7 months.
E2E Networks (India)	Rs.200/hr	Rs.1,44,000	2.8 months	Indian cloud, data sovereignty, affordable for startups.
RunPod H100	Rs.172/hr	Rs.1,23,955	3.3 months	Budget choice. Spot interruptions unreliable for long training runs.
Vast.ai (marketplace)	Rs.162/hr	Rs.1,16,429	3.5 months	Cheapest option. Community nodes - not reliable for production.
Oracle Cloud (Bare Metal)	Rs.108/hr/GPU	Rs.77,850	5.2 months	Cheapest hyperscaler per GPU. Worth for cost-sensitive workloads.

3. Indian Cloud & Datacenter Infrastructure (100% Data Sovereignty)

All providers maintain 100% data sovereignty. Critical for DPDP Act compliance, healthcare, BFSI, and government AI workloads in India.

> **IndiaAI Mission subsidised rates:** Rs.115-150/hr for eligible Indian startups and researchers building foundation AI models. 40% subsidy for other GPU workloads. Apply at indiaai.gov.in.

Provider	Category	GPU Hardware	Hourly (INR)	Monthly (INR)	Data Sovereignty	Quality	Efficiency	IndiaAI	Key Notes
Jio Cloud (Reliance Jio)	Sovereign GPU Cloud	NVIDIA GH200 Grace Hopper	**Rs.200-350/hr**	**Rs.1.4-2.5L**	100% Sovereign - Azure-backed	★★★★☆	★★★★☆	**Yes**	450M Jio users. NVIDIA GH200+DGX Cloud. 2,000 MW DC ambition Jamnagar.
Yotta Shakti Cloud	GPU Cloud (Largest India)	H100/GH200/B200 - 32,768 GPUs	**Rs.115-150 (subsidy) / Rs.250-400**	**Rs.83K-2.9L**	100% Sovereign - NM+Telangana	★★★★★	★★★★★	**Yes**	India's first NVIDIA NCP. NM1 campus 210 MW. Clients: Sarvam AI, Zoho.
Sify CloudInfinit+AI (NASDAQ:SIFY)	DGX-Ready GPU Cloud	H100/H200/L4 - DGX-Ready 130kW/rack	**Rs.200-380/hr**	**Rs.1.4-2.7L**	100% - Chennai/Noida/Mumbai	★★★★☆	★★★★☆	**Yes**	DGX-Ready certified. Pay-per-use. 14 DCs, 407+ MW.
AdaniConneX (Adani+EdgeConneX)	Hyperscale Colocation	Customer-brought GPUs. AI-ready racks.	**Colo pricing (per kW/month)**	**Rs.60K-1.5L/rack/mo**	100% - Chennai/Noida/Vizag	★★★★☆	★★★★☆	**Partner**	Google JV: 1 GW AI campus Visakhapatnam ($15B, 2026-30). Bring-your-own-GPU.
E2E Networks (NSE-listed)	GPU Cloud	NVIDIA Hopper + Quantum-2 InfiniBand	**Rs.150-250/hr (billed per minute)**	**Rs.1.1-1.8L**	100% India DCs - NSE-listed	★★★★☆	★★★★☆	**Yes**	Hosts AI4Bharat, Qure.ai, Invideo AI. Best for Indian startups.
Tata Communications AI Cloud	Enterprise GPU Cloud	NVIDIA Hopper + Blackwell roadmap 2026	**Rs.220-400/hr**	**Rs.1.6-2.9L**	100% India-owned - global network	★★★★☆	★★★★☆	**Yes**	AI Studio, NIM, Omniverse. ~30% of world internet routes.
CDAC / PARAM Rudra / IndiaAI HPC	Govt HPC (Research Only)	PARAM Rudra - H100/A100/AMD MI300X	**Rs.115-150/hr (40-100% subsidy)**	**Rs.83K-1.1L (subsidised)**	100% Govt-owned sovereign	★★★☆☆	★★★☆☆	**Core provider**	NOT commercial. Apply via IndiaAI Mission portal. Research/foundation model only.

4. India AI Ecosystem - Full Entity Classification Directory

Green rows = confirmed DC/cloud infra providers. Amber = hardware/infra vendors.

Entity / Name	Category	DC / Cloud?	Description & Role in AI Ecosystem
AdaniConneX	Hyperscale DC	Yes - DC/Colo	50:50 JV Adani+EdgeConneX. AI-ready DCs Chennai/Noida/Vizag. Google JV 1 GW AI campus Visakhapatnam ($15B, 2026-30).
Jio Cloud	Sovereign GPU Cloud	Yes - GPU Cloud	GPU-as-a-Service on NVIDIA GH200 & DGX Cloud. Fully in India, Azure-backed sovereign infra. IndiaAI Mission empanelled.
Yotta Shakti Cloud	GPU Cloud (Largest)	Yes - GPU Cloud	32,768 H100/GH200/B200 GPUs. India's first NVIDIA NCP. NM1 210 MW campus. Clients: Sarvam AI, Zoho.
Sify CloudInfinit+AI	DGX-Ready GPU Cloud	Yes - GPU Cloud	NVIDIA DGX-Ready (Mumbai/Chennai/Noida, 130kW/rack). Pay-per-use May 2025. IndiaAI empanelled H100/H200/L4.
E2E Networks	GPU Cloud (NSE-listed)	Yes - GPU Cloud	NSE-listed Indian GPU cloud. H100 Hopper+Quantum-2 InfiniBand. Hosts AI4Bharat, Qure.ai, Invideo AI.
Tata Comm. AI Cloud	Enterprise GPU Cloud	Yes - GPU Cloud	NVIDIA Hopper+Blackwell roadmap. AI Studio, NIM, Omniverse. ~30% world internet routes.
AnantRaj Cloud	Hyperscale DC (nascent)	DC (building)	Building hyperscale AI DCs Delhi-NCR/Manesar. MoU with Haryana for Rs.50,000 crore campus. GPU cloud nascent.
CDAC Rudra Server	HPC Server Platform	HPC Node	Made-in-India server (Intel Xeon/AMD EPYC+MI300X). 1300+ units in PARAM supercomputers. NOT rentable - internal HPC.
Schneider Electric	DC Power/Cooling Infra	HW Vendor	UPS, PDUs, cooling, DCIM. GTC 2026 partner in Vera Rubin DSX AI Factory design. Supplier to Yotta/Sify/Adani.
Supermicro	Server Hardware OEM	HW Vendor	AI server chassis (HGX H100/B200). Used by Yotta and Indian cloud builders. NVIDIA partner.
Arista Networks	DC Networking	HW Vendor	400G+ switches for AI cluster spine-leaf fabric. Used in Indian hyperscale DCs.
CommScope	Cabling / Physical Infra	HW Vendor	Structured cabling, fiber management inside DCs. Physical infra enabler.
Corning	Fiber Optics	HW Vendor	World's largest fiber optic maker. DC interconnects and subsea cables.
DDN (Data Direct Networks)	AI Storage	HW Vendor	Parallel storage for AI training pipelines. GTC 2026: storage partner in NVIDIA BlueField-4 STX ecosystem.
AMD Versal/Strix/Vision	Chip/AI Accelerator	Chip Vendor	MI300X in CDAC Rudra-GN. Strix Halo=Ryzen AI Max. IndiaAI accepts AMD MI300X/MI325X from empanelled providers.

Entity / Name	Category	DC / Cloud?	Description & Role in AI Ecosystem
Velankani Group	IT Infra / ODM	**Network vendor**	Indian ODM making white-box network hardware. Supplies DC networking.
Tata (in chips)	Semiconductor Fab	**Chip fab**	Tata Electronics chip manufacturing Dholera. India Semiconductor Mission.

5. NVIDIA Vera CPU & Dynamo 1.0 - GTC March 16, 2026 (Breaking Announcements)

BREAKING - GTC 2026: Vera CPU and Dynamo 1.0 announced/launched at NVIDIA GTC on March 16, 2026. Dynamo 1.0 is available FREE TODAY on existing Blackwell GPUs for up to 7x speedup.

AI Stack Layer	Technology (March 2026)	Description & Function	Relevance to India / On-Prem
Application / Agent Layer	OpenClaw / NemoClaw / NVIDIA Agent Toolkit (GTC 2026)	AI agents, LLM apps, RAG pipelines. NemoClaw = NVIDIA secure enterprise agent OS. One-command install on DGX Spark, RTX PC, or cloud VM.	Available NOW on DGX Spark, TAALAS, cloud. NemoClaw open-source.
Inference Orchestration	NVIDIA Dynamo 1.0 - Open-source 'OS for AI Factories' (GA March 16, 2026)	Disaggregates prefill (GPU) and decode (Groq LPU). Up to 7x inference boost on same Blackwell hardware. KV-cache-aware smart routing. Integrates vLLM, SGLang, TensorRT-LLM, LangChain. Used by AWS/Azure/GCP/CoreWeave/ByteDance/PayPal.	Deploy FREE TODAY on TAALAS or H100 cloud for 7x speedup. India providers (Yotta, Sify) adopting.
Processing / Compute Layer	NVIDIA Vera CPU + Rubin GPU + Groq 3 LPX (Announced GTC 2026)	Vera CPU = NVIDIA custom CPU for agentic AI (data prep, tool execution, code). Pairs with Rubin GPU (heavy compute) + Groq 3 LPX (decode, 35x throughput/megawatt). Vera Rubin NVL72 = rack-scale AI supercomputer. 10x throughput/watt vs Blackwell.	Hardware coming 2026. DGX Spark uses GB10 Grace CPU (predecessor). Monitor Yotta/Sify for Rubin GPU cloud.
Hardware / Silicon Layer	Rubin GPU / Groq 3 LPX / BlueField-4 STX / ConnectX-9 / Spectrum-X 400G+	Rubin GPU handles training+prefill. Groq 3 LPX handles low-latency decode. BlueField-4 manages networking/KV storage. Spectrum-X 400G+ Ethernet scale-out.	Current DGX Spark = Blackwell (predecessor gen). Full Vera Rubin platform 2026.
Datacenter / Power / Cooling	Vera Rubin DSX AI Factory reference design (Schneider, Siemens, Eaton, Vertiv, Cadence)	NVIDIA blueprint for building AI factories. Omniverse DSX digital twin for DC operations. Partners: Schneider Electric, DDN, Dell, VAST Data, HPE.	Indian DCs (Yotta, Sify, AdaniConneX) implementing AI-factory designs per NVIDIA DSX reference.

6. LLM Cost & Load Reduction Techniques

These techniques reduce GPU memory usage, inference time, cloud bills, and power consumption. Applicable to all platforms. Stack multiple techniques for compounding benefits.

Technique	Memory	Speed Gain	Energy	Quality Impact	Tools / Frameworks	Description & Usage
INT4 / Q4 Quantization	4-8x	2-3x	40-60%	1-3% drop	*GPTQ, AWQ, llama.cpp, TensorRT-LLM*	Reduce FP16 to INT4. 70B needs 140 GB -> 35 GB VRAM. Runs on single RTX 4090 or DGX Spark. Most practical technique for deployment.
FP4 Quantization (NVIDIA native)	8-10x	3-5x	50-70%	2-5% drop	*TensorRT-LLM FP4, DGX Spark native*	Blackwell-native FP4 with sparsity. DGX Spark does 1 PFLOP FP4 natively. Jetson Thor also Blackwell FP4. 70B model at FP4 needs ~17.5 GB.
KV Cache + Offloading (Dynamo KVBM)	N/A	+30-60% tput	20-40%	None	*Dynamo 1.0, vLLM PagedAttention, SGLang*	Cache Key-Value pairs. Smart routing to GPUs already holding context. Offload cold cache to CPU RAM/NVMe. BitNet a4.8 supports 3-bit KV cache.
Context Optimisation (Compression)	N/A	2-16x	40-80%	Neutral/improves	*LLMLingua, RAG, NVIDIA Grove*	Inference scales O(n^2) with context. Cutting 32K to 8K context = ~16x less attention compute. Replace long context with RAG retrieval.
Microsoft BitNet b1.58 (1-bit LLM)	Up to 32x	2-6x (CPU)	72-82%	Minimal (2B scale)	*bitnet.cpp (Microsoft open-source)*	Ternary weights {-1,0,+1}. Natively trained. 2B model = 0.4 GB. Runs 100B on single CPU at 5-7 tok/s. 82% energy reduction. MUST use bitnet.cpp.
Speculative Decoding	N/A	2-4x decode	20-30%	None (exact output)	*TensorRT-LLM, SGLang, vLLM*	Small draft model proposes tokens; large model verifies in parallel. Works on DGX Spark and TAALAS. Jetson Thor supports Eagle-based speculative decoding.
Continuous Batching (PagedAttention)	N/A	3-10x GPU util	30-50%	None	*vLLM PagedAttention, SGLang*	Process requests as they arrive. 15% GPU utilization -> 80%+. Essential for TAALAS multi-user production. Standard in all modern inference engines.
NVIDIA Dynamo 1.0 (GA March 2026)	N/A	Up to 7x	30-50%	None	*Dynamo, NIXL, KVBM, Grove (open-source)*	Disaggregated prefill/decode. KV-cache-aware routing. Available NOW for free on Blackwell. AWS, Azure, GCP, CoreWeave, ByteDance all using it.
Model Distillation + LoRA	70-90%	5-15x	70-85%	Task-specific loss	*NVIDIA NIM, Nemo, LoRA/QLoRA*	Small student model mimics large teacher. DeepSeek R1 distilled to 1.5B-7B retains 90%+ reasoning on specific tasks. Great for task-specific APIs.
FULL STACK (INT4+KV+batch+Dynamo)	4-8x	7-20x	60-80%	1-4% drop	*All above combined*	Best production setup. TAALAS running 70B: serve 15-25 concurrent users at production quality. 4x lower cost than equivalent cloud spend.

7. Recommendations by Use Case

Use Case	Recommendation
Indian developers (data sovereignty required)	Yotta Shakti Cloud or E2E Networks. Apply IndiaAI Mission subsidised compute (Rs.115-150/hr). For enterprises: AdaniConneX or Sify DGX-Ready campuses.
Local daily LLM development (no recurring bill)	NVIDIA DGX Spark Rs.4,06,454 one-time. Runs 200B params, full CUDA. Pays off vs Azure in <30 days. Deploy Dynamo 1.0 for 7x free inference boost.
Production multi-user API (India on-premise)	TAALAS 10x RTX 4090 (~Rs.38-43L). Apply INT4+Dynamo 1.0+batching -> serve 15-25 concurrent users. Pays off vs cloud in ~9-12 months.
Edge AI / Robotics / Kiosk / Embedded LLM	NVIDIA Jetson AGX Thor (~Rs.2,16,250). 128 GB, 2070 FP4 TFLOPS, 40-130W configurable. Full CUDA+vLLM. Perfect for India ATMs, kiosks, robots, healthcare edge.
Personal/team AI exploration (Apple ecosystem)	Mac Mini M4 Pro 64GB (Rs.1,69,900) for 30B, or Mac Studio M4 Max 128GB (~Rs.4,19,000) for 70B. Silent, ultra-efficient. NOT for multi-user production. No CUDA.
Cost-sensitive global compute (no India req.)	Oracle Cloud ($1.25/GPU/hr = Rs.108/hr) cheapest hyperscaler. RunPod (Rs.172/hr) and Vast.ai (Rs.162/hr) cheapest overall but NOT production-reliable.
Government / academic research in India	Apply IndiaAI Mission compute portal (indiaai.gov.in). PARAM Rudra / H100 / AMD MI300X at 40-100% subsidy. Sarvam AI received 4,096 H100 GPUs free.
Max inference performance (new hardware 2026+)	Vera Rubin platform coming 2026: 10x throughput/watt vs Blackwell. Deploy Dynamo 1.0 NOW (free, open-source) for 7x speedup on existing hardware.

8. Quick Reference Summary - All Platforms

Platform	Max Model	Price (INR)	Multi-user	Data Sovereignty	CUDA	Power	Best For / Notes
Mac Mini M4 16GB	7B INT4	Rs.57,900	1 user	Full	✗	35 W	Entry personal AI. Silent. No CUDA.
Mac Mini M4 Pro 64GB	30B INT4	Rs.1,69,900	2-3 users	Full	✗	55 W	Best value personal workstation.
Mac Studio M4 Max 128GB	70B INT4	~Rs.4,19,000	3-5 users	Full	✗	100 W	Silent 70B inference. No CUDA.
NVIDIA DGX Spark	200B FP4	Rs.4,06,454	3-8 users	Full	✓	170 W	Best local CUDA dev. Pays off vs Azure <30 days.
NVIDIA Jetson AGX Thor	70B FP8	~Rs.2,16,250	8-16 (small)	Full	✓	80 W	Edge AI champion. Robotics/kiosk/India.
TAALAS - 10x RTX 4090	70-130B INT4	Rs.38-43L	15-25 users	Full	✓	3 kW	Production powerhouse. Needs active/liquid cooling.
Cloud H100 80GB (1x GPU)	70B FP16	Rs.162-604/hr	8-20+ scaled	Varies	✓	DC	Best for burst training. Fully scalable.
Raspberry Pi 5 + AI Kit	1-3B only	~Rs.7,200	Demo only	Full	✗	12 W	Education/prototype only.
Yotta Shakti (India)	All sizes	Rs.115-400/hr	Scalable	100%	✓	DC	Best sovereign GPU cloud India.
Jio Cloud (India)	All sizes	Rs.200-350/hr	Scalable	100%	✓	DC	Sovereign, Azure-tech, 450M users.
E2E Networks (India)	All sizes	Rs.150-250/hr	Scalable	100%	✓	DC	Best for Indian startups/researchers.

PRICING DISCLAIMER All pricing as of March 2026. Cloud GPU prices are highly dynamic - verify current rates before procurement decisions. IndiaAI Mission pricing applies only to eligible Indian entities building foundation AI models. DGX Spark MSRP = $4,699 / Rs.4,06,454 (post Feb 2026 price revision). Monthly cloud cost = hourly rate x 8 hrs x 30 days. 1 USD = Rs.86.5.